

Import the important libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn import metrics
```

Import The Dataset

```
#loading the data from csv file to a Pandas DataFrame
insurance_dataset = pd.read_csv('/content/insurance.csv')

insurance_dataset.head()

{"summary": {"name": "insurance_dataset", "rows": 1338,
 "fields": [{"column": "age", "dtype": "number", "min": 14, "max": 64, "std": 14, "num_unique_values": 47, "samples": [21, 21, 21, 21, 21], "semantic_type": "\\", "description": "\n"}, {"column": "sex", "dtype": "category", "num_unique_values": 2, "samples": ["male", "female"], "semantic_type": "\\", "description": "\n"}, {"column": "bmi", "dtype": "number", "min": 15.96, "max": 53.13, "std": 6.098186911679017, "num_unique_values": 548, "samples": [23.18, 23.18, 23.18, 23.18, 23.18], "semantic_type": "\\", "description": "\n"}, {"column": "children", "dtype": "number", "min": 0, "max": 5, "std": 1, "num_unique_values": 6, "samples": [0, 0, 0, 1, 1, 1], "semantic_type": "\\", "description": "\n"}, {"column": "smoker", "dtype": "category", "num_unique_values": 2, "samples": ["no", "yes"], "semantic_type": "\\", "description": "\n"}], "properties": {}}, {"column": "region", "dtype": "category", "num_unique_values": 4, "samples": ["southeast", "northeast", "midwest", "west"], "semantic_type": "\\", "description": "\n"}]
```

```

\"northeast\"],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n    },\n    {\n      \"column\":\n        \"charges\", \n        \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 12110.011236693994, \n          \"min\": 1121.8739, \n          \"max\": 63770.42801, \n          \"num_unique_values\": 1337, \n          \"samples\": [\n            8688.85885, \n            5708.867\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n        }\n      }\n    }\n  },\n  \"type\": \"dataframe\", \"variable_name\": \"insurance_dataset\"}\n\ninsurance_dataset.shape\n(1338, 7)\n\ninsurance_dataset.info()\n<class 'pandas.core.frame.DataFrame'>\nRangeIndex: 1338 entries, 0 to 1337\nData columns (total 7 columns):\n #   Column   Non-Null Count   Dtype \n--- \n 0   age       1338 non-null    int64 \n 1   sex       1338 non-null    object \n 2   bmi       1338 non-null    float64 \n 3   children  1338 non-null    int64 \n 4   smoker    1338 non-null    object \n 5   region    1338 non-null    object \n 6   charges   1338 non-null    float64 \n dtypes: float64(2), int64(2), object(3)\nmemory usage: 73.3+ KB

```

Categorical Features:

1) Sex 2) Smoker 3) Region

```

insurance_dataset.isnull().sum()\n\nage      0\nsex      0\nbmi      0\nchildren 0\nsmoker   0\nregion   0\ncharges  0\ndtype: int64\n\ninsurance_dataset.describe()\n\n{ \"summary\": {\"name\": \"insurance_dataset\", \"rows\": 8,\n  \"fields\": [\n    {\n      \"column\": \"age\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 

```

```

460.6106090399993,\n          \"min\": 14.049960379216172,\n\"max\": 1338.0,\n          \"num_unique_values\": 8,\n\"samples\": [\n            39.20702541106129,\n              39.0,\n              ],\n              \"semantic_type\": \"\",\n              \"column\":\n              \"bmi\", \n              \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 463.29524977918294,\n                  \"min\": 6.098186911679017,\n                  \"max\": 1338.0,\n                  \"num_unique_values\": 8,\n                  \"samples\": [\n                    30.66339686098655,\n                      30.4,\n                      ],\n                      \"semantic_type\": \"\", \n                      \"column\":\n                      \"children\", \n                      \"properties\": {\n                        \"dtype\": \"number\", \n                        \"std\": 472.5368318870757,\n                          \"min\": 0.0,\n                          \"max\": 1338.0,\n                          \"num_unique_values\": 7,\n                          \"samples\": [\n                            1338.0,\n                              1.0949177877429,\n                            ],\n                            \"semantic_type\": \"\", \n                            \"column\":\n                            \"charges\", \n                            \"properties\": {\n                              \"dtype\": \"number\", \n                              \"std\": 20381.922846226596,\n                                \"min\": 1121.8739,\n                                \"max\": 63770.42801,\n                                \"num_unique_values\": 8,\n                                \"samples\": [\n                                  13270.422265141257,\n                                    9382.033,\n                                  ],\n                                  \"semantic_type\": \"\", \n                                  \"column\":\n                                  \"description\": \"\"\n                                } \n                              } \n                            } \n                          },\n                          \"type\": \"dataframe\"}\n
```

```

# distribution of age value
sns.set()
plt.figure(figsize=(6,6))
sns.distplot(insurance_dataset['age'])
plt.title('Age Distribution')
plt.show()

```

/tmp/ipython-input-3634923312.py:4: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```

sns.distplot(insurance_dataset['age'])

```



```
# Gender column
plt.figure(figsize=(6,6))
sns.countplot(x='sex', data=insurance_dataset)
plt.title('Sex Distribution')
plt.show()
```



```
insurance_dataset['sex'].value_counts()

sex
male      676
female    662
Name: count, dtype: int64


# bmi distribution
plt.figure(figsize=(6,6))
sns.distplot(insurance_dataset['bmi'])
plt.title('BMI Distribution')
plt.show()

/tmp/ipython-input-1916795400.py:3: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn
v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```


```
For a guide to updating your code to use the new functions, please see https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751
```

```
sns.distplot(insurance_dataset['bmi'])
```



Normal BMI Range --> 18.5 to 24.9


```
# children column
plt.figure(figsize=(6,6))
sns.countplot(x='children', data=insurance_dataset)
plt.title('Children')
plt.show()
```



```
insurance_dataset['children'].value_counts()

children
0    574
1    324
2    240
3    157
4     25
5     18
Name: count, dtype: int64


# smoker column
plt.figure(figsize=(6,6))
sns.countplot(x='smoker', data=insurance_dataset)
plt.title('smoker')
plt.show()
```



```
insurance_dataset['smoker'].value_counts()

smoker
no      1064
yes     274
Name: count, dtype: int64

# region column
plt.figure(figsize=(6,6))
sns.countplot(x='region', data=insurance_dataset)
plt.title('region')
plt.show()
```



```
insurance_dataset['region'].value_counts()

region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64

# distribution of charges value
plt.figure(figsize=(6,6))
sns.distplot(insurance_dataset['charges'])
plt.title('Charges Distribution')
plt.show()


/tmp/ipython-input-3971177022.py:3: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn
```

v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(insurance_dataset['charges'])
```



Data Pre-Processing

Encoding the categorical features

```

# encoding sex column
insurance_dataset.replace({'sex':{'male':0,'female':1}}, inplace=True)

3 # encoding 'smoker' column
insurance_dataset.replace({'smoker':{'yes':0,'no':1}}, inplace=True)

# encoding 'region' column
insurance_dataset.replace({'region':
{'southeast':0,'southwest':1,'northeast':2,'northwest':3}},
inplace=True)

/tmp/ipython-input-2871422651.py:2: FutureWarning: Downcasting
behavior in `replace` is deprecated and will be removed in a future
version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
    insurance_dataset.replace({'sex':{'male':0,'female':1}},
inplace=True)
/tmp/ipython-input-2871422651.py:5: FutureWarning: Downcasting
behavior in `replace` is deprecated and will be removed in a future
version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
    insurance_dataset.replace({'smoker':{'yes':0,'no':1}}, inplace=True)
/tmp/ipython-input-2871422651.py:8: FutureWarning: Downcasting
behavior in `replace` is deprecated and will be removed in a future
version. To retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
    insurance_dataset.replace({'region':
{'southeast':0,'southwest':1,'northeast':2,'northwest':3}},
inplace=True)

X = insurance_dataset.drop(columns='charges', axis=1)
Y = insurance_dataset['charges']

print(X)

```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	1
1	18	0	33.770	1	1	0
2	28	0	33.000	3	1	0
3	33	0	22.705	0	1	3
4	32	0	28.880	0	1	3
...
1333	50	0	30.970	3	1	3
1334	18	1	31.920	0	1	2
1335	18	1	36.850	0	1	0
1336	21	1	25.800	0	1	1
1337	61	1	29.070	0	0	3

```
[1338 rows x 6 columns]

print(Y)

0      16884.92400
1      1725.55230
2      4449.46200
3     21984.47061
4     3866.85520
...
1333    10600.54830
1334    2205.98080
1335    1629.83350
1336    2007.94500
1337    29141.36030
Name: charges, Length: 1338, dtype: float64
```

Splitting data into training and test set

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state=2)

print(X.shape, X_train.shape, X_test.shape)

(1338, 6) (1070, 6) (268, 6)
```

Model Training

1. Linear Regression

```
# loading the Linear Regression model
lr_model = LinearRegression()
lr_model.fit(X_train, Y_train)
model = LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None)

# prediction on training data
training_data_prediction = regressor.predict(X_train)

# R squared value
r2_train = metrics.r2_score(Y_train, training_data_prediction)
print('R squared vale : ', r2_train)

R squared vale :  0.751505643411174

# prediction on test data
test_data_prediction_lr = lr_model.predict(X_test)
# R squared value
r2_lr = metrics.r2_score(Y_test, test_data_prediction_lr)
print('R squared vale : ', r2_lr)
```

```
mse = metrics.mean_squared_error(Y_test, test_data_prediction_lr)
rmse = np.sqrt(mse)
print(rmse)

R squared vale : 0.7447273869684076
6191.690842285236
```

Decision tree regressor

```
dt_model = DecisionTreeRegressor(random_state=42)
dt_model.fit(X_train, Y_train)
test_data_prediction_dt = dt_model.predict(X_test)
r2_dt = metrics.r2_score(Y_test, test_data_prediction_dt)
print(f"Decision Tree R² Score: {r2_dt:.4f}")

mse = metrics.mean_squared_error(Y_test, test_data_prediction_dt)
rmse = np.sqrt(mse)
print(rmse)

Decision Tree R² Score: 0.7166
6524.204772451173
```

Random Forest Regressor

```
rf_model = RandomForestRegressor(random_state=42, n_estimators=100)
rf_model.fit(X_train, Y_train)
test_data_prediction_rf = rf_model.predict(X_test)
r2_rf = metrics.r2_score(Y_test, test_data_prediction_rf)
print(f"Random Forest R² Score: {r2_rf:.4f}")

mse = metrics.mean_squared_error(Y_test, test_data_prediction_rf)
rmse = np.sqrt(mse)
print(rmse)

Random Forest R² Score: 0.8379
4933.692230552383
```

XG Boost

```
xgb_model = XGBRegressor(random_state=42, n_estimators=100)
xgb_model.fit(X_train, Y_train)
test_data_prediction_xgb = xgb_model.predict(X_test)
r2_xgb = metrics.r2_score(Y_test, test_data_prediction_xgb)
print(f"XGBoost R² Score: {r2_xgb:.4f}")

mse = metrics.mean_squared_error(Y_test, test_data_prediction_xg)
rmse = np.sqrt(mse)
print(rmse)
```

```

XGBoost R2 Score: 0.8144
5279.090073315875

results = {
    "Linear Regression": {"R2 Score": 0.7447, "RMSE": 6191.69},
    "Decision Tree": {"R2 Score": 0.7166, "RMSE": 6524.60 },
    "Random Forest": {"R2 Score": 0.8379, "RMSE": 4933.69},
    "XGBoost": {"R2 Score": 0.8144, "RMSE": 5279.09}
}

# Convert to DataFrame
results_df = pd.DataFrame(results).T # .T to make models as rows
results_df = results_df.round(4)      # Round values for clean display

print(results_df)

          R2 Score      RMSE
Linear Regression  0.7447  6191.69
Decision Tree     0.7166  6524.60
Random Forest     0.8379  4933.69
XGBoost           0.8144  5279.09

```

Prediction of Medical Insurance on new input data

1. Linear Regression

```

input_data = (31,1,25.74,0,1,0)

# changing input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = lr_model.predict(input_data_reshaped)
print(prediction)

print('The insurance cost is USD ', prediction[0])

[3760.0805765]
The insurance cost is USD  3760.080576496057

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
validation.py:2739: UserWarning: X does not have valid feature names,
but LinearRegression was fitted with feature names
warnings.warn(

```

1. Decision Tree

```

input_data = (31,1,25.74,0,1,0)

# changing input_data to a numpy array

```

```
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = dt_model.predict(input_data_reshaped)
print(prediction)

print('The insurance cost is USD ', prediction[0])

[3756.6216]
The insurance cost is USD 3756.6216

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
validation.py:2739: UserWarning: X does not have valid feature names,
but DecisionTreeRegressor was fitted with feature names
    warnings.warn(
```

Random Forest

```
input_data = (31,1,25.74,0,1,0)

# changing input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = rf_model.predict(input_data_reshaped)
print(prediction)

print('The insurance cost is USD ', prediction[0])

[3729.6420035]
The insurance cost is USD 3729.6420035000065

/usr/local/lib/python3.11/dist-packages/sklearn/utils/
validation.py:2739: UserWarning: X does not have valid feature names,
but RandomForestRegressor was fitted with feature names
    warnings.warn(
```

XG Boost

```
input_data = (31,1,25.74,0,1,0)

# changing input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)
```

```
prediction = xgb_model.predict(input_data_reshaped)
print(prediction)

print('The insurance cost is USD ', prediction[0])

[3409.0215]
The insurance cost is USD 3409.0215
```